

# Chapter 12 - Optimisation

- ▶ Newton's method for optimisation
- ▶ The golden-section method
- ▶ Multivariate optimisation
- ▶ Steepest ascent
- ▶ Newton's method in higher dimensions
- ▶ A curve-fitting example

## 12 - Introduction: 1D Optimisation

Suppose  $f : \mathbb{R} \rightarrow \mathbb{R}$  with continuous first and second derivatives:

- ▶  $f$  has a *global maximum* at  $x^*$  if  $f(x) \leq f(x^*)$  for all  $x$ .
- ▶  $f$  has a *local maximum* at  $x^*$  if  $f(x) \leq f(x^*)$  for all  $x$  in a neighbourhood of  $x^*$ .

Conditions for local maximum:

- ▶ Necessary condition for  $x^*$  to be a local maximum is  $f'(x^*) = 0$  and  $f''(x^*) \leq 0$ .
- ▶ Sufficient condition is  $f'(x^*) = 0$  and  $f''(x^*) < 0$ .

A *local search* technique generates a sequence of points,  $x(0)$ ,  $x(1)$ ,  $x(2)$ ,  $\dots$ , which (hopefully) converge to a local maximum of  $f$ .

- ▶ Given a prospective solution  $x(n)$ , we look for the next prospective solution  $x(n+1)$  in some neighbourhood of  $x(n)$ .

## 12 - Introduction: 1D Optimisation (cont.)

- ▶ Because they never consider the whole space of possible solutions, local search techniques can only ever be guaranteed to find local maxima.
- ▶ Let  $x^*$  be a local maximum of  $f$ . Supposing  $x(n) \rightarrow x^*$  as  $n \rightarrow \infty$ , we need *stopping criteria* to decide when to stop searching:
  - ▶  $|x(n) - x(n-1)| \leq \varepsilon$ ;
  - ▶  $|f(x(n)) - f(x(n-1))| \leq \varepsilon$ ;
  - ▶  $|f'(x(n))| \leq \varepsilon$ .
- ▶ Local search techniques may not converge at all.
  - ▶ For example if  $f$  is unbounded then  $x(n) \rightarrow \infty$ .
  - ▶ Usually specify a maximum number of iterations  $n_{\max}$ .

## 12.1 - Newton's method for optimisation

If  $f : [a, b] \rightarrow \mathbb{R}$  has a continuous derivative  $f'$ , then the maximum of  $f$  is the maximum of

- ▶  $f(a)$ ,  $f(b)$ , and
- ▶  $f(x_1), \dots, f(x_n)$ , where  $x_1, \dots, x_n$  are the roots of  $f'$ .

If we apply the Newton-Raphson method for root-finding to  $f'$ , we get the Newton method for optimising  $f$ :

$$x(n+1) = x(n) - \frac{f'(x(n))}{f''(x(n))}.$$

## 12.1 - Newton's method for optimisation

```
newton <- function(f3, x0, tol = 1e-9, n.max = 100) {  
  # Newton's method for optimisation, starting at x0  
  # f3 is a function that given x returns the vector  
  # (f(x), f'(x), f''(x)), for some f  
  x <- x0  
  f3.x <- f3(x)  
  n <- 0  
  while ((abs(f3.x[2]) > tol) & (n < n.max)) {  
    x <- x - f3.x[2]/f3.x[3]  
    f3.x <- f3(x)  
    n <- n + 1  
  }  
  if (n == n.max) { cat('newton failed to converge\n')}  
  else { return(x)}  
}  
newton(gamma.2.3, 0.25)
```

## 12.1 - Newton's method for optimisation (cont.)

- ▶ When the Newton algorithm converges, we can end up with a minimum or a maximum since all such stationary points satisfy  $f'(x^*) = 0$ .
- ▶ Because we are searching for a point  $x^*$  such that  $f'(x^*) = 0$ , we will use  $|f'(x(n))| < \varepsilon$  as our stopping condition.
- ▶ Provided  $x(0)$  is close to  $x^*$ ,  $x(n) \rightarrow x^*$  quickly, as  $n \rightarrow \infty$ .
- ▶ We revisit Newton's method later, on a higher plane (that is, in higher dimensions).

## 12.2 - The golden-section method


Let  $f : \mathbb{R} \rightarrow \mathbb{R}$  be a continuous function.

The golden-section method works in one dimension only, but does not need  $f'$ .

The golden-section method is similar to the root-bracketing technique for root-finding.

- ▶ To determine if we have a local maximum we need three points: if  $x_l < x_m < x_r$  and  $f(x_l) \leq f(x_m)$  and  $f(x_m) \leq f(x_r)$  then there must be a local maximum in the interval  $[a, b]$ .

## 12.2 - The golden-section method





## 12.2 - The golden-section method: Algorithm

How to choose  $y$ ?:

- ▶ Suppose that  $(x_m, x_r)$  is the larger interval
- ▶ let  $a = x_m - x_l$ ,  $b = x_r - x_m$ , and  $c = y - x_m$ .
- ▶ The golden-section algorithm chooses  $y$  so that the ratio of the lengths of the larger to the smaller interval stays the same at each iteration.
- ▶ That is, if the new bracketing interval is  $[x_l, y]$  then

$$\frac{a}{c} = \frac{b}{a}$$

while if the new bracketing interval is  $[x_m, x_r]$  then

$$\frac{b-c}{c} = \frac{b}{a}.$$

## 12.2 - The golden-section method: Algorithm

- ▶ Put  $\rho = b/a$  then solving these for  $c$  we get

$$\rho^2 - \rho - 1 = 0 \quad \text{so} \quad \rho = \frac{1 + \sqrt{5}}{2}$$

which is the famous golden ratio. We also get  $a = b - c$ , so  $c = b/(1 + \rho)$  and thus  $y = x_m + c = x_m + (x_r - x_m)/(1 + \rho)$ .

- ▶ The length ratio of the new interval to the old is either  $b/(a + b)$  or  $(a + c)/(a + b)$ , which both work out as  $\rho/(1 + \rho)$ .
- ▶ An analogous argument applies if  $(x_l, x_m)$  is the larger interval.

The argument above shows that if we start with  $x_m$  chosen so that the ratio  $(x_r - x_m)/(x_m - x_l) = \rho$  or  $1/\rho$ , then at each iteration the width of the bracketing interval is reduced by a factor of  $\rho/(1 + \rho)$  and so must eventually go to zero.

## 12.2 - The golden-section method: Algorithm

**Golden-section method 2** Start with  $x_l < x_m < x_r$  such that  $f(x_l) \leq f(x_m)$  and  $f(x_r) \leq f(x_m)$

1 if  $x_r - x_l \leq \varepsilon$  then stop

2 if  $x_r - x_m > x_m - x_l$  then do 2a otherwise do 2b

2a let  $y = x_m + (x_r - x_m)/(1 + \rho)$

if  $f(y) \geq f(x_m)$  then put  $x_l = x_m$  and  $x_m = y$  otherwise put  $x_r = y$

2b let  $y = x_m - (x_m - x_l)/(1 + \rho)$

if  $f(y) \geq f(x_m)$  then put  $x_r = x_m$  and  $x_m = y$  otherwise put  $x_l = y$

3 go back to step 1

---

```
gsection <- function(ftn, x.l, x.r, x.m, tol = 1e-9) {  
  # applies the golden-section algorithm to maximise ftn  
  # we assume that ftn is a function of a single variable  
  # and that  $x.l < x.m < x.r$  and  $ftn(x.l), ftn(x.r) \leq ftn(x.m)$   
  #  
  # the algorithm iteratively refines  $x.l$ ,  $x.r$ , and  $x.m$  and terminates  
  # when  $x.r - x.l \leq tol$ , then returns  $x.m$   
  
  # golden ratio plus one  
  gr1 <- 1 + (1 + sqrt(5))/2  
  
  return(x.m)  
}  
  
f <- function(x) ifelse (x==0, 0, abs(x)*log(abs(x)/2)*exp(-abs(x)))  
curve(f, -10, 10, n = 501)  
gsection(f, 1, 5, 2)  
gsection(f, -1, 1, .1)
```

---

## 12.3 - Multivariate Optimisation

Let  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  and suppose that all of the first- and second-order partial derivatives of  $f$  exist and are continuous everywhere.

- ▶ We write  $\mathbf{x} = (x_1, \dots, x_d)^T$  for an element of  $\mathbb{R}^d$  and
- ▶  $\mathbf{e}_i$  for the  $i$ -th co-ordinate vector:  $\mathbf{x} = x_1 \mathbf{e}_1 + \dots + x_d \mathbf{e}_d$ .
- ▶ The  $i$ -th partial derivative at  $\mathbf{x}$  will be denoted  $f_i(\mathbf{x}) = \partial f(\mathbf{x}) / \partial x_i$  and we define the *gradient*

$$\nabla f(\mathbf{x}) = (f_1(\mathbf{x}), \dots, f_d(\mathbf{x}))^T$$

and the *Hessian*

$$\mathbf{H}(\mathbf{x}) = \begin{pmatrix} \frac{\partial^2 f(\mathbf{x})}{\partial x_1 \partial x_1} & \dots & \frac{\partial^2 f(\mathbf{x})}{\partial x_1 \partial x_d} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f(\mathbf{x})}{\partial x_d \partial x_1} & \dots & \frac{\partial^2 f(\mathbf{x})}{\partial x_d \partial x_d} \end{pmatrix}.$$

## 12.3 - Multivariate Optimisation

- ▶ The slope at  $\mathbf{x}$  in direction  $\mathbf{v} \neq \mathbf{0}$  is given by

$$\mathbf{v}^T \nabla f(\mathbf{x}) / \|\mathbf{v}\|,$$

where  $\|\mathbf{v}\| = \sqrt{v_1^2 + \dots + v_d^2}$  is the Euclidean norm.

- ▶ The curvature at  $\mathbf{x}$  in direction  $\mathbf{v}$  is given by

$$\mathbf{v}^T \mathbf{H}(\mathbf{x}) \mathbf{v} / \|\mathbf{v}\|^2.$$

- ▶  $f$  has a local maximum at  $\mathbf{x}$  if for all  $\varepsilon > 0$  small enough,  $f(\mathbf{x} + \varepsilon \mathbf{e}_i) \leq f(\mathbf{x})$  for  $i = 1, \dots, d$ .

## 12.3 - Multivariate Optimisation

- ▶  $f$  has a local maximum at  $\mathbf{x}$  if for all  $\varepsilon > 0$  small enough,  
 $f(\mathbf{x} + \varepsilon \mathbf{e}_i) \leq f(\mathbf{x})$  for  $i = 1, \dots, d$ .
- ▶ A necessary (but not sufficient) condition for a local maximum at  $\mathbf{x}$  is:
  - ▶  $\nabla f(\mathbf{x}) = \mathbf{0} = (0, \dots, 0)^T$  and for all  $\mathbf{v} \neq \mathbf{0}$
  - ▶ the curvature at  $\mathbf{x}$  in direction  $\mathbf{v}$  is  $\leq 0$  (we say that the Hessian is *negative semi-definite*).
- ▶ A sufficient (but not necessary) condition for  $f$  to have a local maximum at  $\mathbf{x}$  is that:
  - ▶  $\nabla f(\mathbf{x}) = \mathbf{0}$  and
  - ▶ the curvature in all directions is  $< 0$  (in which case we say that the Hessian  $\mathbf{H}(\mathbf{x})$  is *negative-definite*).

## 12.3 - Multivariate Optimisation

- ▶ As in one dimension, we will use iterative local search techniques to find local maxima.
- ▶ Define  $\|\mathbf{x}\|_\infty = \max_i |x_i|$  (the  $L_\infty$  norm).
- ▶ In higher dimensions we use stopping conditions that are combinations of the following:
  - $\|\mathbf{x}(n) - \mathbf{x}(n-1)\|_\infty \leq \varepsilon$ ;
  - $|f(\mathbf{x}(n)) - f(\mathbf{x}(n-1))| \leq \varepsilon$ ;
  - $\|\nabla f(\mathbf{x}(n))\|_\infty \leq \varepsilon$ .
- ▶ To guard against non-convergence, we should also specify a maximum number of iterations  $n_{\max}$ , then stop when  $n = n_{\max}$ .



## 12.4 - Steepest Ascent

Put  $\mathbf{x}(n+1) = \mathbf{x}(n) + \alpha \mathbf{v}$ , where  $\alpha$  is a positive scalar and the direction  $\mathbf{v}$  is the direction with largest slope.

At point  $\mathbf{x}$ , the direction with largest slope is  $\nabla f(\mathbf{x})$ .

Thus, the steepest method has the form

$$\mathbf{x}(n+1) = \mathbf{x}(n) + \alpha \nabla f(\mathbf{x}(n)),$$

for some  $\alpha \geq 0$ .

## 12.4 - Steepest Ascent

Given  $\mathbf{x}(n+1) = \mathbf{x}(n) + \alpha \nabla f(\mathbf{x}(n))$ , we choose  $\alpha \geq 0$  to maximise

$$g(\alpha) = f(\mathbf{x}(n) + \alpha \nabla f(\mathbf{x}(n))).$$

- ▶ If  $\alpha = 0$  then we have reached a local maximum.
- ▶ If  $\alpha > 0$  then  $f(\mathbf{x}(n+1)) > f(\mathbf{x}(n))$ .

If  $f$  is bounded above then, because  $f(\mathbf{x}(n+1)) \geq f(\mathbf{x}(n))$ , the sequence  $\{f(\mathbf{x}(n))\}_{n=1}^{\infty}$  must converge. This suggests that we can use the stopping condition:

- ▶  $f(\mathbf{x}(n)) - f(\mathbf{x}(n-1)) \leq \varepsilon$ , for some small tolerance  $\varepsilon$ .

It can be shown that if  $f$  is bounded and  $\nabla f$  is 'well behaved', then the sequence  $\{\mathbf{x}(n)\}_{n=1}^{\infty}$  will converge to a local maximum.

## 12.4 - Steepest Ascent

```
ascent <- function(f, grad.f, x0, tol = 1e-9, n.max = 100) {  
  # steepest ascent algorithm  
  # find a local max of f starting at x0  
  # function grad.f is the gradient of f  
  
  x.old <- x0  
  x <- line.search(f, x0, grad.f(x0))  
  n <- 1  
  while ((f(x) - f(x.old) > tol) & (n < n.max)) {  
    x.old <- x  
    x <- line.search(f, x, grad.f(x))  
    n <- n + 1  
  }  
  return(x)  
}
```

## 12.4.1 - Line Search


To maximise  $g(\alpha) = f(\mathbf{x}(n) + \alpha \nabla f(\mathbf{x}(n)))$  over  $\alpha \geq 0$ , golden-section algorithm will be used.

We require three initial points  $\alpha_l < \alpha_m < \alpha_r$  such that  $g(\alpha_m) \geq g(\alpha_l)$  and  $g(\alpha_m) \geq g(\alpha_r)$ :

- ▶ Put  $\alpha_l = 0$ .
- ▶ In theory, if  $\|\nabla f(\mathbf{x}(n))\| > 0$  then  $g'(0) > 0$  and thus there must be some  $\varepsilon > 0$  such that  $g(\varepsilon) > g(0)$ , so we can put  $\alpha_m = \varepsilon$ .
- ▶ Unfortunately there is not even a theoretical guarantee that a suitable  $\alpha_r$  exists, because we may have  $g$  increasing over the whole interval  $[0, \infty)$ .


Hence, we specify a *maximum step size*  $\alpha_{\max}$  and if we cannot find  $\alpha_r \leq \alpha_{\max}$  such that  $g(\alpha_r) \leq g(\alpha_m)$ , we just return  $\alpha_{\max}$ .

## 12.4.1 - Line Search



## 12.4.1 - Line Search

$$f(x,y)=\sin(x^2/2-y^2/4)*\cos(2*x-\exp(y))$$



## 12.5 Newton's method in higher dimensions

The basis of the method is a second-order Taylor expansion of  $f$ . For any  $\mathbf{x}$  and  $\mathbf{y}$  close together we have

$$f(\mathbf{y}) \approx f(\mathbf{x}) + (\mathbf{y} - \mathbf{x})^T \nabla f(\mathbf{x}) + \frac{1}{2}(\mathbf{y} - \mathbf{x})^T \mathbf{H}(\mathbf{x})(\mathbf{y} - \mathbf{x}).$$

Taking partial derivatives w.r.t.  $\mathbf{y}$  we get

$$\nabla f(\mathbf{y}) \approx \nabla f(\mathbf{x}) + \mathbf{H}(\mathbf{x})(\mathbf{y} - \mathbf{x}).$$

If  $\mathbf{y}$  is a local maximum then  $\nabla f(\mathbf{y}) = \mathbf{0}$  and, solving the equation above, we get

$$\mathbf{y} = \mathbf{x} - \mathbf{H}(\mathbf{x})^{-1} \nabla f(\mathbf{x}).$$

## 12.5 - Newton's method in higher dimensions

Suppose  $\mathbf{x}(n)$  is our current estimate, then we would like our next estimate  $\mathbf{x}(n+1)$  to be a local maximum (at least approximately) ...

### Newton's algorithm

$$\mathbf{x}(n+1) = \mathbf{x}(n) - \mathbf{H}(\mathbf{x}(n))^{-1} \nabla f(\mathbf{x}(n)).$$

- ▶ Clearly if  $\mathbf{H}(\mathbf{x}(n))$  is singular (has no inverse), then Newton's method breaks down.
- ▶ Even if  $\mathbf{H}(\mathbf{x}(n))$  is non-singular at each step, Newton's method may not converge.
- ▶ Despite this, if  $f$  has a local maximum at  $\mathbf{x}^*$ ,  $f$  is 'nicely behaved' near  $\mathbf{x}^*$ , and if our initial point  $\mathbf{x}(0)$  is 'close enough' to  $\mathbf{x}^*$ , then Newton's method will converge to  $\mathbf{x}^*$  quickly.



## 12.5 - Newton's method in higher dimensions


- ▶ In implementing Newton's method we will assume that we have some function  $\mathbf{f3}$  that takes argument  $\mathbf{x}$  and returns a list containing  $f(\mathbf{x})$ ,  $\nabla f(\mathbf{x})$ , and  $\mathbf{H}(\mathbf{x})$ .
- ▶ For our stopping condition we will use  $\|\nabla f(\mathbf{x}(n))\|_{\infty} \leq \varepsilon$ .

---

```
newton <- function(f3, x0, tol = 1e-9, n.max = 100) {  
  # f3 returns the list {f(x), grad f(x), Hessian f(x)}, for some f  
  x <- x0  
  f3.x <- f3(x)  
  n <- 0  
  while ((max(abs(f3.x[[2]])) > tol) & (n < n.max)) {  
    x <- x - solve(f3.x[[3]], f3.x[[2]])  
    f3.x <- f3(x)  
    n <- n + 1  
  }  
  if (n == n.max) { cat('newton failed to converge\n') }  
  else { return(x) }
```

## 12.5 - Newton's method in higher dimensions

$$f(x,y)=\sin(x^2/2-y^2/4)*\cos(2*x-\exp(y))$$



## 12.5 - Newton's method in higher dimensions

- ▶ Newton's method needs to calculate the gradient and Hessian.
- ▶ Steepest ascent only requires the gradient, but sometimes even this can be difficult.

If  $\nabla f$  is unavailable then there are two approaches we can take:

- ▶ The first assumes that even if we don't know what they are,  $\mathbf{H}$  and/or  $\nabla f$  do exist, in which case we can try and estimate them.
- ▶ The second approach is to use an optimisation method that does not require the gradient.
  - ▶ Such approaches tend to be relatively slow, but relatively reliable. In one dimension the golden-section algorithm is an example of a derivative-free approach.
  - ▶ In higher dimensions there is an algorithm due to Nelder & Mead, which is well accepted and again is derivative-free.

## 12.7 - Curve fitting

Suppose we have observations  $(x_1, y_1), \dots, (x_n, y_n)$  and we want to find a function  $f$  such that  $y_i \approx f(x_i)$  for  $i = 1, \dots, n$ .

Further suppose that  $f$  can be *parameterised* by some vector of parameters  $\theta = (\theta_1, \dots, \theta_d)^T$ . For example, if we restrict  $f$  to be a quadratic then it has the form  $f(x) = ax^2 + bx + c$ , in which case  $\theta = (a, b, c)^T$ .

We write  $f(x; \theta)$  for  $f(x)$  to emphasise the dependence on  $\theta$ .

The problem of finding the parameter  $\theta^*$ , such that the fitted points  $\hat{y}_i = f(x_i; \theta^*)$  are 'closest' to the observations  $y_i$ , is called *curve fitting*.

To measure how close the fitted points are to the observed points, we use a *loss function*. Two popular choices are the sum of squares

$$L_2(\theta) = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$


and the sum of absolute differences

$$L_1(\theta) = \sum_{i=1}^n |y_i - \hat{y}_i|.$$

Note that we consider a loss function to be a function of  $\theta$ , rather than a function of  $\mathbf{y}$ , because we are interested in how the loss changes as we change  $\theta$ .

Given a loss function  $L$ , we choose  $\theta^*$  to be that  $\theta$  that minimises  $L(\theta)$ .

**Tree 1.3.11**



## 12.7 - Curve fitting

---

```
richards <- function(t, theta)
  theta[1]*(1 - exp(-theta[2]*t))^theta[3]
```

```
loss.L2 <- function(theta, age, vol)
  sum((vol - richards(age, theta))^2)
```

```
loss.L1 <- function(theta, age, vol)
  sum(abs(vol - richards(age, theta)))
```

```
trees <- read.csv("../data/trees.csv")
tree <- trees[trees$ID=="1.3.11", 2:3]
```

```
theta0 <- c(1000, 0.1, 3)
theta.L2 <- optim(theta0, loss.L2, age=tree$Age, vol=tree$Vol)
theta.L1 <- optim(theta0, loss.L1, age=tree$Age, vol=tree$Vol)
```

---

## 12.7 - Curve fitting

---

```
plot(tree$Age, tree$Vol, type="p", xlab="Age", ylab="Volume",  
      main="Tree 1.3.11")  
lines(tree$Age, richards(tree$Age, theta.L2$par), col="blue")  
lines(tree$Age, richards(tree$Age, theta.L1$par), col="blue", lty=2)  
  
attach(tree)  
fit <- nls(Vol ~ a*(1 - exp(-b*Age))^c, start=list(a=1000, b=0.01,  
            c=3))  
summary(fit)
```

---